

RCB allocation vs naive scheduling — $L = 8$ layers, 2-qubit chain, per-qubit dephasing

